



TRAVERSE RESEARCH

Bindless Rendering in D3D12

NEXT SLIDE





BINDLESS RENDERING

OVERVIEW

- 01 Shader model pre-6.6 & 6.6 bindless.
- 02 Rendering resource handles
- 03 Descriptor management.
- 04 Bindless shader code.
- 05 Bonus



Bindless Rendering

MOTIVATION

- 01 Significantly reduce API overhead.
- 02 Reduce rendering pipeline complexity.
- 03 Descriptor staging is a thing of the past.
- 04 Makes DXR 1.0 / 1.1 a lot easier to work with.
- 05 More lightweight rendering backend.



TRADITIONAL "BINDFULL" PROBLEMS

ROOT SIGNATURES

Each pipeline has a unique root signature.

Sometimes uses shader reflection.

A lot of management code and/or factories to make this all work.

DESCRIPTOR STAGING

Staging descriptors from a big heap on the cpu to smaller gpu heaps.

A lot of management code.

Difficult to track.

EXCESSIVE API CALLS

`CopyDescriptors` call per unique draw/dispatch.

`SetCompute/
SetGraphicsRootDescriptor` per unique draw/dispatch

The point of d3d12 is to have lower CPU overhead, let's make it count!





BINDLESS EVERYTHING

BINDLESS RENDERING

Setting up bindless heaps

- A single shader visible descriptor heap for all CBV_SRV_UAV descriptors and optionally SAMPLER.
- Non-shader visible heaps still required for RTV and DSV.
- CBV_SRV_UAV descriptor heap can go up to 1,000,000+ descriptors
- SAMPLER heap is limited (2048) on most hardware.
- Make sure to check out hardware tiers. We assume hardware tier 3 for our framework “Breda”.



SHADER MODEL 6.5 OR LOWER

BINDLESS RENDERING

A “global” root signature.

- A single root signature for all shaders.
- Assign register space per resource type.
- Set up root constants.
- Set up static samplers.

```
let mut descriptor_ranges: Vec<D3D12_DESCRIPTOR_RANGE1> = vec![];

// Texture2D, repeat for all resource types.
descriptor_ranges.push(d3d12::D3D12_DESCRIPTOR_RANGE1 {
    RangeType: d3d12::D3D12_DESCRIPTOR_RANGE_TYPE_SRV,
    NumDescriptors: bindless_descriptor_count,
    BaseShaderRegister: 0,
    RegisterSpace: BindlessTableType::Texture2d.space_index() as u32,
    Flags: d3d12::D3D12_DESCRIPTOR_RANGE_FLAG_DESCRIPTOR_VOLATILE
        | d3d12::D3D12_DESCRIPTOR_RANGE_FLAG_DATA_VOLATILE,
    OffsetInDescriptorsFromTableStart: 0,
});
```

SHADER MODEL 6.6+

```
let mut desc: D3D12_VERSIONED_ROOT_SIGNATURE_DESC = d3d12::D3D12_VERSIONED_ROOT_SIGNATURE_DESC {
    Version: d3d12::D3D_ROOT_SIGNATURE_VERSION_1_1,
    ..Default::default()
};

let desc_1_1: &mut D3D12_ROOT_SIGNATURE_DESC1 = unsafe { desc.u.Desc_1_1_mut() };
desc_1_1.NumParameters = bindless_params.len() as u32;
desc_1_1.pParameters = bindless_params.as_ptr();
desc_1_1.NumStaticSamplers = static_samplers.len() as u32;
desc_1_1.pStaticSamplers = static_samplers.as_ptr().cast();
desc_1_1.Flags = d3d12::D3D12_ROOT_SIGNATURE_FLAG_CBV_SRV_UAV_HEAP_DIRECTLY_INDEXED
    | d3d12::D3D12_ROOT_SIGNATURE_FLAG_SAMPLER_HEAP_DIRECTLY_INDEXED;
```



BINDLESS RENDERING

Rendering resource handles

- A *RenderResourceHandle* maps 1:1 to an index in our bindless descriptor heap.
- *RenderResourceHandle* is simply a uint32 but we are not using all 32 bits, unused bits can be used for validation purposes.
- *RenderResourceHandles* are exclusively created during resource creation.
- What about SRV & UAV as they are separate descriptors?

```
impl RenderResourceHandle {
    pub fn new(version: u8, tag: RenderResourceTag, index: u32) -> Self {
        let version: u32 = version as u32;
        let tag: u32 = tag as u32;
        let index: u32 = index as u32;

        // version wraps around, it's just to make sure invalid resources don't get another version
        assert!(version < 64);
        assert!(tag < 8);
        assert!(index < (1 << 23));

        Self(version << 26 | tag << 23 | index)
    }
}
```

HLSL

```
struct RenderResourceHandle {
    uint handle;

    uint read_index() { return this.handle & ((1 << 23) - 1); }
    bool is_valid() { return this.handle != ~0; }

#ifdef VK_BINDLESS
    uint write_index() { return this.handle & ((1 << 23) - 1); }
#else
    uint write_index() {
        uint uav_idx = this.handle + 1;
        return uav_idx & ((1 << 23) - 1);
    }
#endif
};
```



BINDLESS RENDERING

Rendering resource handles

- Introduction of *RenderResourceHandlePair*
- SRV & UAV descriptors are assumed at index { [n], [n+1] }
- Each (sub)resource has a *RenderResourceHandlePair*

IMPLEMENTATION

```
pub struct RenderResourceHandlePair {
    pub srv: RenderResourceHandle,
    pub uav: RenderResourceHandle,
}

impl RenderResourceHandlePair {
    pub fn new(version: u8, tag: RenderResourceTag, descriptor_idx: u32) -> Self {
        Self::new_from(srv: RenderResourceHandle::new(version, tag, index: descriptor_idx))
    }

    /// Create an srv-uav pair from the `srv`, with `uav` on the next index
    pub fn new_from(srv: RenderResourceHandle) -> Self {
        RenderResourceHandlePair {
            srv,
            uav: srv.with_index(srv.index() + 1),
        }
    }

    /// Return start of the handle pair at index N (in this case SRV) where N + 1 contains the UAV.
    pub fn handle(self) -> RenderResourceHandle {
        self.srv
    }
}
```



BINDLESS RENDERING

Rendering resource handles

- Each resource has a *RenderResourceHandlePair*, additionally Render targets/Depth stencils may have additional RTV / DSV handles in their according heap types.
- Create the views once during resource creation.

VIEWS CREATED DURING RESOURCE INIT

```
let descriptor_pair: RenderResourceHandlePair = descriptor_pool.allocate_buffer_handle_pair();

device.CreateShaderResourceView(
    pResource: resource,
    pDesc: &srv_desc,
    DestDescriptor: descriptor_pool.make_cpu_handle(
        heap_type: d3d12::D3D12_DESCRIPTOR_HEAP_TYPE_CBV_SRV_UAV,
        bindless_descriptor: descriptor_pair.srv,
    ),
);

device.CreateUnorderedAccessView(
    pResource: resource,
    pCounterResource: std::ptr::null_mut(),
    pDesc: &uav_desc,
    DestDescriptor: descriptor_pool.make_cpu_handle(
        heap_type: d3d12::D3D12_DESCRIPTOR_HEAP_TYPE_CBV_SRV_UAV,
        bindless_descriptor: descriptor_pair.uav,
    ),
);
```



BINDLESS RENDERING

Resource handle management

- Upon dropping resources, recycle the descriptor.
- Upon allocating RenderResourceHandlePairs, check for recycled handles.
- RenderResourceHandlePairs are recycled in FIFO order.
- Permanent fragmentation may occur if FIFO order is not respected.
- Only recycle descriptors when they are no longer in use!

```
pub(crate) fn retire_handle(&self, handle: RenderResourceHandle) {
    self.available_recycled_descriptors: Arc<Mutex<VecDeque<RenderResourceHandle>>>
        .lock(): Result<MutexGuard<VecDeque<...>>, ...>
        .unwrap(): MutexGuard<VecDeque<RenderResourceHandle>>
        .push_back(handle);
}
```

```
pub(crate) fn retire_cbv_srv_uav_handle(&self, handle: RenderResourceHandlePair) {
    // Uav handles are implicitly retired as they are paired with SRV's
    self.pool[d3d12::D3D12_DESCRIPTOR_HEAP_TYPE_CBV_SRV_UAV as usize].retire_handle(handle.srv);
}
```

```
pub fn allocate_descriptor_pair(&self, tag: RenderResourceTag) -> RenderResourceHandlePair {
    self.available_recycled_descriptors: Arc<Mutex<VecDeque<RenderResourceHandle>>>
        .lock(): Result<MutexGuard<VecDeque<...>>, ...>
        .unwrap(): MutexGuard<VecDeque<RenderResourceHandle>>
        .pop_front(): Option<RenderResourceHandle>
        .map_or_else(
            default: || {
                let descriptor_idx: u32 = self.increment_descriptor_pair();
                RenderResourceHandlePair::new(version: 0, tag, descriptor_idx)
            },
            f: |recycled_handle: RenderResourceHandle| {
                RenderResourceHandlePair::new_from(
                    srv: recycled_handle.bump_version_and_update_tag(tag),
                )
            },
        )
}
```



BINDLESS RENDERING

Preparing for bindless rendering

- 01 Set descriptor heaps.
- 02 Set graphics/compute root signatures.
- 03 (sm6.5) bind descriptor tables.

BIND HEAPS

```
let mut handles: [*mut ID3D12DescriptorHeap; 2] = [  
    descriptor_pool.pool[d3d12::D3D12_DESCRIPTOR_HEAP_TYPE_CBV_SRV_UAV as usize]: DescriptorHeap  
        .handle: NonNull<ID3D12DescriptorHeap>  
        .as_ptr(),  
    descriptor_pool.pool[d3d12::D3D12_DESCRIPTOR_HEAP_TYPE_SAMPLER as usize]: DescriptorHeap  
        .handle: NonNull<ID3D12DescriptorHeap>  
        .as_ptr(),  
];  
  
direct_entry: Arc<CommandBufferPoolEntry>  
    .cmd_list: *mut ID3D12GraphicsCommandList5  
    .as_ref(): Option<&ID3D12GraphicsCommandList5>  
    .unwrap(): &ID3D12GraphicsCommandList5  
    .SetDescriptorHeaps(NumDescriptorHeaps: handles.len() as u32, ppDescriptorHeaps: handles.as_mut_ptr());
```

SET DESCRIPTOR TABLES (SM6.5)

```
direct_entry: Arc<CommandBufferPoolEntry>  
    .cmd_list: *mut ID3D12GraphicsCommandList5  
    .as_ref(): Option<&ID3D12GraphicsCommandList5>  
    .unwrap(): &ID3D12GraphicsCommandList5  
    .SetComputeRootDescriptorTable(  
        RootParameterIndex: 0,  
        BaseDescriptor: descriptor_pool.pool  
            [d3d12::D3D12_DESCRIPTOR_HEAP_TYPE_CBV_SRV_UAV as usize]: DescriptorHeap  
            .unwrap_gpu_heap_start(),  
    );
```



Bindless Rendering

DESCRIPTOR SETS

- 01 Create a buffer containing RenderResourceHandles.
- 02 The created buffer has a RenderResourceHandle.
- 03 Set root constant with the handle of the created buffer.

BUILD BUFFER WITH RENDER HANDLES

```
let set: Arc<dyn DescriptorSet> = DescriptorSetBuilder::persistent(&memcpy_pipeline);
.set.read(index: 0, resource: &input): DescriptorSetBuilder
.set.write(index: 1, resource: &output): DescriptorSetBuilder
.set.build(device, &mut dma);
```

SET HANDLE WITH PUSH CONSTANT

```
fn update_graphics_descriptor_set_handle_push_const(
    &mut self,
    descriptor_set: &Arc<dyn DescriptorSet>,
) {
    let set: &DescriptorSet = descriptor_set.downcast_ref::<Dx12DescriptorSet>().unwrap();
    let binding_handle: RenderResourceHandle = set.buffer.resource_handle();

    unsafe {
        self.cmd().SetGraphicsRoot32BitConstants(
            RootParameterIndex: 1,
            Num32BitValuesToSet: 1,
            pSrcData: (&binding_handle.as_raw() as *const u32).cast(),
            DestOffsetIn32BitValues: PushConstantSlots::Bindings.index() as u32,
        );
    }
}
```



BINDLESS RENDERING

Shader model 6.5 bindless shaders

04 Load bindings in shader.

05 Dynamically index in descriptor heap.

ROOT CONSTANTS

```
struct BindingsOffset {  
    RenderResourceHandle bindingsOffset;  
    uint userData0;  
    uint userData1;  
    uint userData2;  
};
```

DECLARATION OF ALL RESOURCES

```
ConstantBuffer<BindingsOffset> g_bindingsOffset : register(b0, space10);  
  
SamplerState g_samplerMinMagMipPointWrap : register(s0, space0);  
SamplerState g_samplerMinMagMipPointClamp : register(s1, space0);  
SamplerState g_samplerMinMagMipLinearWrap : register(s2, space0);  
SamplerState g_samplerMinMagMipLinearClamp : register(s3, space0);  
  
#define BindlessTexture2DDecl(T) \br/>    Texture2D<T> g_texture2d[] : register(t0, space0)  
#define BindlessTextureCubeDecl(T) \br/>    TextureCube<T> g_textureCube[] : register(t0, space1)
```

DYNAMIC INDEXING 6.5

```
#define texture_2d(handle) \br/>    g_texture2d[NonUniformResourceIndex(handle.read_index())]
```

DYNAMIC INDEXING 6.6

```
#define texture_2d(handle) \br/>    ResourceDescriptorHeap[NonUniformResourceIndex(handle.read_index())]
```



BINDLESS RENDERING

Bindless shader example

- pre-sm6.6 approach is forced to use declarations and functions for dynamic indexing.
- We only use (RW)ByteAddressBuffers for buffer types, templated load and store helps a lot.
- Still not ideal, sm6.6 template approach is ideal.

```
#include "breda-internal::bindless.hlsl"

struct Bindings {
    RenderResourceHandle inputTexture;
    RenderResourceHandle constants;
    RenderResourceHandle outputTexture;
};

struct Constants {
    float foo;
    int bar;
    float3 multiplier;
};

BindlessByteAddressBufferDecl();
BindlessRWTexture2DDecl(float4);
BindlessTexture2DDecl(float4);

[numthreads(8, 8, 1)] void main(uint2 pos
                                : SV_DispatchThreadID) {
    Bindings bnd =
        byte_buffer(g_bindingsOffset.bindingsOffset).Load<Bindings>(0);

    Constants constants = byte_buffer(bnd.constants).Load<Constants>(0);

    Texture2D<float4> some_texture = texture_2d(bnd.inputTexture);
    RWTexture2D<float4> output_target = rw_texture_2d(bnd.outputTexture);

    float4 result = float4(1,0,0,1) * constants.multiplier;

    output_target[pos] = result;
}
```



BONUS SLIDE

Future directions with templating

- More specialized RenderHandles with the use of templates.

```
struct RWBindlessBuffer {
    RenderResourceHandle handle;

    template<typename T>
    T Load(uint index) {
        StructuredBuffer<T> buffer = ResourceDescriptorHeap[self.handle.read_index()];
        return buffer[index];
    }

    template<typename T>
    void Store(uint index, T value) {
        RWStructuredBuffer<T> buffer = ResourceDescriptorHeap[self.handle.write_index()];
        buffer[index] = value;
    }

    template<typename T>
    T LoadBytes(uint offset) {
        ByteAddressBuffer buffer = ResourceDescriptorHeap[self.handle.read_index()];
        return buffer.Load<T>(offset);
    }

    template<typename T>
    void StoreBytes(uint offset, T value) {
        RWByteAddressBuffer buffer = ResourceDescriptorHeap[self.handle.write_index()];
        buffer.Store<T>(offset, value);
    }
};
```



A wide-angle landscape photograph featuring a two-lane asphalt road that curves through a mountainous region. The mountains are rugged and partially covered in snow, with some patches of brown vegetation visible. The sky is filled with soft, colorful clouds in shades of orange, yellow, and blue, suggesting a sunset or sunrise. The overall mood is serene and majestic.

THANK YOU